

# Weight Lift Exercise Quality Prediction

*Jozef M H Dassen*

*Sunday, May 17, 2015*

This is a Machine Learning project attempting to decide the proper way of weight lifting exercises from a set of body movement sensors.

## Data Set

The data set consists of processed sensor measurements. The processing was done using sliding windows on time series measurements of sensors. See <http://groupware.les.inf.puc-rio.br/har> (<http://groupware.les.inf.puc-rio.br/har>) (section on the Weight Lifting Exercise) for details. Each data row contains 152 measurements and is classified to a CLASS A,B,C,D or E indicating the quality of the exercise execution. Many of the columns however contain NA or DIV/0 values. These are removed. This will leave us with 52 usable measurements. No doubt some of these will be correlated since we use different sensor types for measuring the same movements.

We calculate the correlation matrix for the 52 measurements and visualize it with corrplot.



##	freqRatio	percentUnique	zeroVar	nzv
## roll_belt	1.018868	7.94933392	FALSE	FALSE
## pitch_belt	1.043796	12.29526097	FALSE	FALSE
## yaw_belt	1.055710	13.01594235	FALSE	FALSE
## total_accel_belt	1.057173	0.21110868	FALSE	FALSE
## gyros_belt_x	1.085624	0.96090850	FALSE	FALSE
## gyros_belt_y	1.155125	0.47317464	FALSE	FALSE
## gyros_belt_z	1.093496	1.18657640	FALSE	FALSE
## accel_belt_x	1.043165	1.15745796	FALSE	FALSE
## accel_belt_y	1.138837	1.01186576	FALSE	FALSE
## accel_belt_z	1.100000	2.10380724	FALSE	FALSE
## magnet_belt_x	1.075099	2.20572177	FALSE	FALSE
## magnet_belt_y	1.149660	2.09652763	FALSE	FALSE
## magnet_belt_z	1.024465	3.13751183	FALSE	FALSE
## roll_arm	45.673077	17.49290238	FALSE	FALSE
## pitch_arm	72.000000	20.39018709	FALSE	FALSE
## yaw_arm	32.534247	19.31280483	FALSE	FALSE
## total_accel_arm	1.017350	0.46589503	FALSE	FALSE
## gyros_arm_x	1.008242	4.58615418	FALSE	FALSE
## gyros_arm_y	1.491573	2.70801485	FALSE	FALSE
## gyros_arm_z	1.072000	1.67431026	FALSE	FALSE
## accel_arm_x	1.095652	5.58346073	FALSE	FALSE
## accel_arm_y	1.193333	3.84363398	FALSE	FALSE
## accel_arm_z	1.076923	5.57618112	FALSE	FALSE
## magnet_arm_x	1.000000	9.57268690	FALSE	FALSE
## magnet_arm_y	1.076923	6.20222756	FALSE	FALSE
## magnet_arm_z	1.092105	9.14318993	FALSE	FALSE
## roll_dumbbell	1.085106	86.57639950	FALSE	FALSE
## pitch_dumbbell	2.382979	84.29060202	FALSE	FALSE
## yaw_dumbbell	1.068182	86.04498799	FALSE	FALSE
## total_accel_dumbbell	1.127409	0.31302322	FALSE	FALSE
## gyros_dumbbell_x	1.034965	1.68886948	FALSE	FALSE
## gyros_dumbbell_y	1.184652	1.96549465	FALSE	FALSE
## gyros_dumbbell_z	1.041763	1.42680352	FALSE	FALSE
## accel_dumbbell_x	1.013393	3.02103807	FALSE	FALSE
## accel_dumbbell_y	1.029070	3.29766325	FALSE	FALSE
## accel_dumbbell_z	1.229814	2.85360705	FALSE	FALSE
## magnet_dumbbell_x	1.126050	7.78918250	FALSE	FALSE
## magnet_dumbbell_y	1.154472	5.99839849	FALSE	FALSE
## magnet_dumbbell_z	1.146341	4.82638131	FALSE	FALSE
## roll_forearm	11.639485	13.67838684	FALSE	FALSE
## pitch_forearm	66.146341	19.27640678	FALSE	FALSE
## yaw_forearm	15.761628	12.86307054	FALSE	FALSE
## total_accel_forearm	1.161476	0.49501347	FALSE	FALSE
## gyros_forearm_x	1.094183	2.03829075	FALSE	FALSE
## gyros_forearm_y	1.014652	5.26315789	FALSE	FALSE
## gyros_forearm_z	1.155425	2.14020528	FALSE	FALSE
## accel_forearm_x	1.123077	5.66353643	FALSE	FALSE
## accel_forearm_y	1.025974	7.06122152	FALSE	FALSE
## accel_forearm_z	1.119266	4.10569993	FALSE	FALSE
## magnet_forearm_x	1.017857	10.61367111	FALSE	FALSE
## magnet_forearm_y	1.031250	13.27072869	FALSE	FALSE
## magnet_forearm_z	1.022727	11.69105336	FALSE	FALSE
## classe	1.469526	0.03639805	FALSE	FALSE

It turns out that all the variables have sufficiently large variation in measurement to be usable.

We now look at which columns can be removed from the predictors based on the correlation matrix. We use the `findCorrelation` function with 75% as the cut off point. This means we find predictors that are more than 75% correlated with other predictors already in use.

```
findCorrelation(corMatrix, cutoff=0.75)
```

```
## [1] 10 1 9 4 36 8 2 37 35 38 21 34 23 25 12 48 19 46 45 31
```

```
names(train[findCorrelation(corMatrix, cutoff=0.75)])
```

```
## [1] "accel_belt_z"      "roll_belt"         "accel_belt_y"
## [4] "total_accel_belt"  "accel_dumbbell_z"  "accel_belt_x"
## [7] "pitch_belt"       "magnet_dumbbell_x" "accel_dumbbell_y"
## [10] "magnet_dumbbell_y" "accel_arm_x"       "accel_dumbbell_x"
## [13] "accel_arm_z"      "magnet_arm_y"     "magnet_belt_y"
## [16] "accel_forearm_y"  "gyros_arm_y"      "gyros_forearm_z"
## [19] "gyros_forearm_y"  "gyros_dumbbell_x"
```

## Model selection

The problem is a multi level classification problem. Linear or logistic regression is not suitable. A decision tree algorithm seems more suitable and we choose the random forest algorithm. Random forest is reputed to be robust and recommended as a first cut (see <https://www.kaggle.com/wiki/RandomForests> (<https://www.kaggle.com/wiki/RandomForests>)).

The main issue we are facing that needs to be decided is which predictors to choose. We have total of 52 predictors but there are correlations between them as shown by the Correlation matrix. We will attempt to reduce the number of predictors. We will remove the highly correlated predictors found above from the predictors. But we will, as an alternative, also use PCA to reduce the predictors. To investigate the effect of the various methods we will run four models:

1 Remove High Correlation	we remove the highly correlated predictors
2 PCA with threshold 0.80	we use PCA to find predictors keeping 80% of the variance.
3 PCA with threshold 0.95	we use PCA to find predictors keeping 95% of the variance.
4 Use all predictors	we use all 52 predictors

We will evaluate accuracy and training time for each of the 4 cases.

## Cross Validation and Accuracy

In order to determine the Out of Sample error rate of the trained model we need to have a dataset independent of the training data. Therefore we split the original data set into a training data set and a cross validation data set as follows:

```
inTrain <- createDataPartition(y=training$classe,p=0.7, list=FALSE)
train    <- training[inTrain,]
crossval <- training[-inTrain,]
```

We also notice that the range of values for the different predictors can be different by factor of up to 10,000. Therefore we preprocess each data row and scale to center the values (using mean and standard deviation). This preprocessing will be applied to both training and cross validation data sets (and later the test data set as well).

The *train* dataset will be used for training the model and the *crossval* will be used for predicting outcome using the trained model. We then compare the predicted outcome with the actual outcome in a confusion matrix as follows:

```
preProc <- preProcess(train[-lastCol], method = c("scale","center"))
trainPC <- predict(preProc, train1[-lastCol])
crossvalPC <- predict(preProc, crossval[-lastCol])
modelFit <- train(train$classe ~ .,method="rf",data=trainPC)
CM <- confusionMatrix(predict(modelFit,newdata=crossvalPC),crossval$classe)
```

The confusion matrix will give us the Accuracy of the prediction of the outcome on the cross validation data set by the model trained.

As explained above we run 4 different models. Two of the models will use PCA, principle components as predictors. In those cases the preprocessing changes to:

```
preProc <- preProcess(train[-lastCol], method = c("scale","center","pca"), thresh = 0.95)
```

Where *thresh* is the percentage of retained variance required.

We give the full code for the first model, which removes the highly correlated predictors, below. The model and Confusion matrix output are shown.

```
train1 <- train[,-highCorr]
cv1 <- crossval[,-highCorr]
lastCol <- dim(train1)[2]

preProc1 <- preProcess(train1[-lastCol], method = c("scale","center"))
trainPC1 <- predict(preProc1, train1[-lastCol])
cvPC1 <- predict(preProc1, cv1[-lastCol])
tic1=proc.time()[3]
modelFit1 <- train(train1$classe ~ .,method="rf",data=trainPC1)
toc1=proc.time()[3] - tic1
CM1 <- confusionMatrix(predict(modelFit1,newdata=cvPC1),cv1$classe)
modelFit1
CM1
```

```
## Random Forest
##
## 13737 samples
##    30 predictor
##    5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
##
## Summary of sample sizes: 13737, 13737, 13737, 13737, 13737, 13737, ...
##
## Resampling results across tuning parameters:
##
##  mtry  Accuracy   Kappa     Accuracy SD   Kappa SD
##    2    0.9861119 0.9824375  0.001544131   0.001946190
##   16    0.9848733 0.9808723  0.001123780   0.001414139
##   31    0.9755716 0.9691116  0.003222698   0.004064923
##
## Accuracy was used to select the optimal model using  the largest value.
## The final value used for the model was mtry = 2.
```

#### ## Confusion Matrix and Statistics

```
##
##           Reference
## Prediction   A    B    C    D    E
##           A 1674    4    0    0    0
##           B    0 1133    2    0    0
##           C    0    2 1022    9    0
##           D    0    0    2  955    1
##           E    0    0    0    0 1081
```

#### ## Overall Statistics

```
##
##           Accuracy : 0.9966
##           95% CI : (0.9948, 0.9979)
##    No Information Rate : 0.2845
##    P-Value [Acc > NIR] : < 2.2e-16
```

```
##           Kappa : 0.9957
```

```
## McNemar's Test P-Value : NA
```

#### ## Statistics by Class:

```
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      1.0000   0.9947   0.9961   0.9907   0.9991
## Specificity      0.9991   0.9996   0.9977   0.9994   1.0000
## Pos Pred Value   0.9976   0.9982   0.9894   0.9969   1.0000
## Neg Pred Value    1.0000   0.9987   0.9992   0.9982   0.9998
## Prevalence       0.2845   0.1935   0.1743   0.1638   0.1839
## Detection Rate   0.2845   0.1925   0.1737   0.1623   0.1837
## Detection Prevalence 0.2851   0.1929   0.1755   0.1628   0.1837
## Balanced Accuracy 0.9995   0.9972   0.9969   0.9950   0.9995
```

We see that the model uses 31 predictors and the resulting Accuracy is 0.99. We timed the execution using tic/toc variables and find that this model need 2205 seconds to train.

As explained we train 3 additional models, two using PCA and one with the full set of predictors. The result is shown in the table below.

Case	Predictors	Accuracy	Elapsed Time
Remove High Correlation	31	.99	2204.7
PCA with threshold 0.80	12	.96	1082.2
PCA with threshold 0.95	25	.98	1967.8
Use all predictors	52	.99	3446.8

It appears that removing the highly correlated predictors gives the same accuracy as using all predictors. Execution time improves by 36%, which roughly corresponds to the reduction in the number of predictors (40%). The Accuracy of the PCA runs depends on the threshold requested. A threshold value of 95% gives good accuracy but not as good as the models with high correlation predictors removed. The model with high correlation predictors removed seems to be the best compromise. We get a good reduction in computation time without losing accuracy.

## Prediction of Test cases

We will now proceed to predict the outcome of the test data set using the trained models. The test data set needs to be prepared by applying the same preprocessing as was used in the training. We need to take care obviously to apply the correct preprocessing (since we have 4 models). We show the case for our preferred model, the High-Correlation-Removed model:

```
test    <- testing[,-highCorr]
lastCol <- dim(train)[2]
testPC  <- predict(preProc, test[-lastCol])
predict <- predict(modelFit, newdata=testPC)
```

Applying the same process for each of the other 3 models, we get the predictions as shown in the table below. Only the model with PCA at threshold 80% shows different predictions for test sample 3 and 11.

##	Remove High Corr.	PCA threshold 0.80	PCA threshold 0.95	All predictors
## 1	B	B	B	B
## 2	A	A	A	A
## 3	B	A	B	B
## 4	A	A	A	A
## 5	A	A	A	A
## 6	E	E	E	E
## 7	D	D	D	D
## 8	B	B	B	B
## 9	A	A	A	A
## 10	A	A	A	A
## 11	B	A	B	B
## 12	C	C	C	C
## 13	B	B	B	B
## 14	A	A	A	A
## 15	E	E	E	E
## 16	E	E	E	E
## 17	A	A	A	A
## 18	B	B	B	B
## 19	B	B	B	B
## 20	B	B	B	B

## Conclusion

Random forest is an easy ML model to use in the given case. Training time is the only issue to be addressed by attempting to reduce the number of predictors used in the model. Simply removing highly correlated predictors is as effective, if not more so, than using PCA with high variance retention. The trained model works well on the cross validation set.